

# Review of Basic Statistical Procedures

# Basic Statistical Inference

Let  $X_1, \dots, X_n$  be a random sample with pdf  $f(\theta)$ , where either  $f$  or  $\theta$  or both may be unknown.

## Inference:

- Estimation
- Hypothesis Testing

## Approaches:

- Frequentist
- Bayesian

# Parameter Estimation

- The Maximum Likelihood Method

Let  $X_1, \dots, X_n$  be a random sample from  $f(x, \theta)$ . The *likelihood function* of  $\theta$  is defined to be

$$L(\theta) = \prod_i^n f(x_i; \theta)$$

Under regularity conditions,  $\hat{\theta}$ , the MLE, is approximately normal, with mean  $\theta$  and variance  $\frac{1}{nI(\theta)}$ , where

$$I(\theta) = E \left[ \frac{\partial}{\partial \theta} \log f(x, \theta) \right]^2$$

- The Method of Moments

Define the  $k$ 'th *moment* of  $f(x, \theta)$  as

$$\mu_k = E[X^k].$$

The corresponding *sample moments* are given as

$$\hat{\mu}_k = \frac{1}{n} \sum_{i=1}^n X_i^k$$

Let  $\theta_j = h_j(\mu_1, \dots, \mu_s)$  be parameters of interest.

The method of moments estimates of  $\theta_j$

$$\hat{\theta}_j = h_j(\hat{\mu}_1, \dots, \hat{\mu}_s)$$

## *Unbiasedness*

An estimator  $\hat{\theta}$  is said to be unbiased if  $E[\hat{\theta}] = \theta$ .

An estimator is said to be uniformly minimum variance unbiased (UMVU) if it has the smallest variance in the class of unbiased estimators uniformly in  $\theta$ .

## *Consistency*

An estimator  $\hat{\theta}_n$  of  $\theta$  is said to be *consistent* if it approaches  $\theta$  for large  $n$ . More formally, for any  $\epsilon > 0$ ,

$$Pr[|\hat{\theta}_n - \theta| > \epsilon] \rightarrow 0, \quad n \rightarrow \infty$$

Method of moments estimators and most mle's are consistent.

## *Efficiency*

Let  $\tilde{\theta}$  and  $\hat{\theta}$  be possible estimates of  $\theta$ . The *efficiency* of  $\hat{\theta}$  relative to  $\tilde{\theta}$  is defined to be

$$e(\hat{\theta}, \tilde{\theta}) = \frac{\text{Var}(\tilde{\theta})}{\text{Var}(\hat{\theta})}$$

*Cramer-Rao Inequality.* Let  $X_1, \dots, X_n$  be i.i.d. with p.d.f.  $f(x; \theta)$ . Let  $T = T(X_1, \dots, X_n)$  be an unbiased estimator of  $\theta$ . Then under regularity conditions

$$\text{Var}(T) \geq \frac{1}{nI(\theta)}$$

An unbiased estimator whose variance attains the lower bound is said to be *efficient*.

## *Sufficiency*

A statistic  $T = T(X_1, \dots, X_n)$  is said to be sufficient for  $\theta$  if the conditional distribution of the sample given  $T = t$  is independent of  $\theta$ , i.e.,  $T$  contains all the information about  $\theta$ .

A necessary and sufficient condition for  $T$  to be sufficient for  $\theta$  is given by the *factorization theorem*

$$f(x_1, \dots, x_n; \theta) = g(T, \theta)h(x_1, \dots, x_n)$$

If there is a sufficient statistic for  $\theta$  then the mle is also a function of  $T$



*Rao-Blackwell Theorem.* Let  $\hat{\theta}$  be an estimator of  $\theta$ , with  $E[\hat{\theta}^2] < \infty$ , and put  $\tilde{\theta} = E[\hat{\theta} \mid T]$ . Then

$$E(\tilde{\theta} - \theta)^2 \leq E(\hat{\theta} - \theta)^2$$

with strict equality when  $\tilde{\theta} \neq \hat{\theta}$ .

In *complete* families of distributions, if there is an unbiased estimator that is a function of the sufficient statistic, then one should look no further, i.e., that estimator is the unique unbiased minimum variance estimator.

*Lehmann-Scheffe* theorem.

# Hypothesis Testing

Suppose we are interested in determining whether  $\theta \in \Theta_1$  versus  $\theta \in \Theta_2$ , where  $\Theta_1$  and  $\Theta_2$  are disjoint subsets of the parameter space.

Suppose  $X_1, \dots, X_n$  is a random sample from  $f(x, \theta)$ , the hypotheses may be written as

$$H_0 : \theta \in \Theta_0$$

$$H_1 : \theta \in \Theta_1.$$

- Type I error occurs when  $H_0$  is falsely rejected  
Type II: When  $H_0$  is not rejected when it should be.

The probability of Type I error is known as *level of significance*.

Suppose  $H_0$  specifies that the pdf of the data is  $f_0(x)$ , and  $H_1$  claims  $f_1(x)$ .

*Neyman-Pearson Lemma.* Let  $\phi^*$  be a test that accepts when

$$\frac{f_0(x)}{f_1(x)} > c$$

with significance level  $\alpha^*$ . Let  $\phi$  be another test of level  $\alpha \leq \alpha^*$ . Then the power of  $\phi^*$  is greater than or equal to that of  $\phi$ .

For testing composite hypotheses,  $H_0 : \theta \in \Theta_0$  vs  $H_1 : \theta \in \Theta_1$ , the *generalized likelihood ratio tests* may be constructed:

$$\Lambda = \frac{\sup_{\theta \in \Theta_0} [L(\theta)]}{\sup_{\theta \in \Theta_1} [L(\theta)]}$$

where  $\Theta_0$  and  $\Theta_1$  are partitions of the parameter space  $\Theta$ .

In application, one uses the following formulation

$$\Lambda = \frac{\sup_{\theta \in \Theta_0} [L(\theta)]}{\sup_{\theta \in \Theta} [L(\theta)]}$$

The likelihood ratio test rejects for small values of  $\Lambda$ , or large values of  $-\log \Lambda$ .

Given a test statistic  $T = T(X_1, \dots, X_n; \theta_o)$ , and the observed value  $t_{obs}$ , the *p-value* is defined as the probability of observing a value of  $T$  as extreme as or more extreme than  $t_{obs}$ . Small p-values are evidence against  $H_o$ .



## Confidence Intervals

Let  $L = L(X_1, \dots, X_n)$  and  $U = U(X_1, \dots, X_n)$  be statistics, such that  $Pr[L \leq \theta \leq U] = 1 - \alpha$ . Then the interval  $[L, U]$  is called a  $100(1 - \alpha)\%$  confidence interval for  $\theta$ .

# Assumptions of Classical Statistical Inference

- Independence
- Normality (or other specified distribution)

## Validation of Assumptions:

- Histograms, qqnorm, etc., for normality
- Box-plots, stem-and-leaf diagrams for outlier detection
- Simple time series plot for serial correlations: including trends and cycles.

When distribution of data skewed, confidence intervals tend to be large, on the average, and p-values may be inflated.

Nonparametric methods do not make explicit assumptions about underlying distributions.

Most commonly used nonparametric methods are based on rank transformations of the observations.

Example: Cedar-apple rust data:

- Disease that affects apple trees, with symptom, rust-colored spots on apple leaves. Study conducted to study cause.
- In the first year of experiment the number of affected leaves on a random sample of 8 trees was counted on Variety 1.
- Normally, the average number of affected leaves is at most 25.

Data:

Tree	Count
1	38
2	10
3	84
4	36
5	50
6	35
7	73
8	48

Does the data suggest the apple trees are affected?

# Standard Tests for Location

## *One-Sample Problem*

Let  $X_1, \dots, X_n$  be a random sample from a distribution with mean  $\mu$  and variance  $\sigma^2$ . We wish to test the hypothesis

$$H_0 : \mu = \mu_0$$

vs.

$$H_1 : \mu > \mu_0$$

When the distribution is normal, and  $\sigma^2$  unknown (the usual case), a procedure that is commonly used is the Student's t-test, given by

$$T = \frac{\sqrt{n}(\bar{X} - \mu_0)}{S}$$

which under the null has a  $t_{n-1}$  distribution.

Example: Cedar-apple rust data:

```
> variety1
```

```
[1] 38 10 84 36 50 35 73 48
```

```
> t.test(variety1,mu=25)
```

One Sample t-test

data: variety1

t = 2.651, df = 7, p-value = 0.03289

alternative hypothesis: true mean is not equal to 25

95 percent confidence interval:

27.34963 66.15037

sample estimates:

mean of x: 46.75

If EDA shows that the data is non-normal, then the *Wilcoxon signed-rank* test may be used.

- Let  $D_i = X_i - \mu_0$
- Rank  $|D_1|, \dots, |D_n|$
- Let  $T_+$  and  $T_-$  be the sums of the ranks assigned to the positive and negative  $d_i$ 's, respectively. Clearly, if  $H_0$  is true,  $T_+ \approx T_-$ .

Further, under  $H_0$ ,

$$E(T_+) = \frac{n(n+1)}{4}$$

and

$$Var(T_+) = \frac{n(n+1)(2n+1)}{24}$$



- The Wilcoxon signed-rank test rejects for large values of  $T_+$ , or when

$$Z = \frac{T_+ - E[T_+ | H_0]}{\sqrt{Var(T_+ | H_0)}}$$

exceeds a standard normal critical point.

- The test requires the underlying distribution to be symmetric.

### Example: Cedar-apple rust data (cont'd):

- Disease that affects apple trees, with symptom, rust-colored spots on apple leaves. Study conducted to study cause.
- In the first year of experiment the number of affected leaves on a random sample of 8 trees from Variety 1 and a random sample of 7 trees from Variety 2 counted.

Data:

Variety 1		Variety 2	
Tree	Count	Tree	Count
1	38	1	27
2	10	2	28
3	84	3	57
4	36	4	66
5	50	5	77
6	35	6	49
7	73	7	62
8	48		

Does the data suggest there is a difference between the two varieties in the mean number of affected leaves?

### *Two-Sample Problem*

Let  $X_1, \dots, X_n$ , and  $Y_1, \dots, Y_m$ , be independent random samples from distributions with respective means,  $\mu_1$  and  $\mu_2$ , and variances,  $\sigma_1^2$  and  $\sigma_2^2$ .

We are interested in the hypothesis

$$H_0 : \mu_1 - \mu_2 = \Delta_0$$

vs

$$H_1 : \mu_1 - \mu_2 > \Delta_0$$

When both distributions are normal, with unknown, but equal, variance  $\sigma^2$ , the two-sample t-test is given by

$$T = \frac{(\bar{X} - \bar{Y}) - \Delta_0}{\sqrt{S_p^2(\frac{1}{n} + \frac{1}{m})}}$$

where the pooled variance  $S_p^2$  is defined as

$$S_p^2 = \frac{(n-1)S_1^2 + (m-1)S_2^2}{n+m-2}$$

It is easy to show that  $T$  has a  $t_{n+m-2}$  distribution.

Example: Cedar-apple rust data (cont'd):

```
> variety2 <- c(27,28,57,66,77,49,62)
```

```
> t.test(variety1,variety2,var.equal=T)
```

Two Sample t-test

data: variety1 and variety2

$t = -0.501$ ,  $df = 13$ ,  $p\text{-value} = 0.6247$

alternative hypothesis:

true difference in means is not equal to 0

95 percent confidence interval:

-29.40555 18.33412

sample estimates:

mean of x mean of y

46.75000 52.28571

Example: Cedar-apple rust data (cont'd):

```
> var.test(variety1,variety2)
```

F test to compare two variances

data: variety1 and variety2

F = 1.499, num df = 7, denom df = 6,

p-value = 0.638

alternative hypothesis:

true ratio of variances is not equal to 1

95 percent confidence interval:

0.2631926 7.6728062

sample estimates:

ratio of variances

1.499006

When the variances are not equal, an approximate test is *Welch's modified two-sample t test*, given by

$$T = \frac{(\bar{X} - \bar{Y}) - \Delta_0}{S_w}$$

where

$$S_w^2 = \frac{S_1^2}{n} + \frac{S_2^2}{m}$$

T has an approximate  $t_\nu$  distribution, with

$$\nu \approx \left[ \frac{c^2}{n-1} + \frac{(1-c)^2}{m-1} \right]^{-1}$$

where

$$c = \frac{S_1^2}{nS_w^2}$$

Example: Cedar-apple rust data (cont'd):

```
> t.test(variety1,variety2)
```

Welch Two Sample t-test

data: variety1 and variety2

t = -0.5082, df = 12.956,

p-value = 0.6198

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-29.07417 18.00275

sample estimates:

mean of x mean of y


46.75000 52.28571

## Example: Cedar-apple rust data (cont'd):

```
> variety2 <- c(27,28,57,66,77,49,62)
```

```
> variety3 <- c(27,28,57,66,77,49,6200)
```

```
> boxplot(variety2,variety3)
```





## Example: Cedar-apple rust data (cont'd):

```
t.test(variety1,variety2,var.equal=T)
```

Two Sample t-test

data: variety1 and variety2

t = -0.501, df = 13,

p-value = 0.6247

alternative hypothesis:

true difference in means is not  
equal to 0

95 percent confidence interval:

-29.40555 18.33412

sample estimates:

mean of x mean of y

46.75000 52.28571

```
> variety3 <- c(27,28,57,66,77,49,620)
```

```
> t.test(variety1,variety3,var.equal=T)
```

Two Sample t-test

data: variety1 and variety3

t = -1.1151, df = 13, p-value = 0.285

alternative hypothesis: true difference in  
means is not equal to 0

95 percent confidence interval:

-250.4077 79.9077

sample estimates:

mean of x mean of y

46.75 132.00

# Wilcoxon Rank-Sum Test

- Combine the two samples, and rank the observations
- Let  $T_1$  be the sum of the ranks for the observations in the 1st group, and  $T_2$  for the second
- The Wilcoxon rank-sum test is then based on  $T_W$ , which has an asymptotic standard normal distribution,

$$T_W = \frac{T_1 - E(T_1 \mid H_o)}{\sqrt{\text{Var}(T_1 \mid H_o)}}$$

where

$$E(T_1 \mid H_o) = \frac{n(n + m + 1)}{2}$$

and

$$\text{Var}(T_1 \mid H_o) = \frac{nm(n + m + 1)}{12}$$

In some applications, the Mann-Whitney (U) form of the Wilcoxon rank-sum test is used, due to the relative ease of computing critical values using the latter. The two are related by the equation

$$U_1 = T_1 - \frac{n(n+1)}{2}$$

- The test is most appropriate when the populations have the same shape and differ only in location (e.g., same dispersion) However, the distributions do not have to be symmetric.

## Example: Cedar-apple rust data (cont'd):

```
> wilcox.test(variety1,variety2)
```

Wilcoxon rank sum test

data: variety1 and variety2

W = 24, p-value = 0.6943

alternative hypothesis: true mu  
is not equal to 0

```
> variety3 <-  
  c(27,28,57,66,77,49,620)
```

```
> wilcox.test(variety1,variety3)
```

Wilcoxon rank sum test

data: variety1 and variety3

W = 22, p-value = 0.5358

alternative hypothesis: true mu  
is not equal to 0

### Example: Cedar-apple rust data (cont'd):

- Disease that affects apple trees, with symptom, rust-colored spots on apple leaves. Study conducted to study cause.
- In the first year of experiment the number of affected leaves on a random sample of 8 trees from Variety 1
- All red cedar trees within 100 yards of the orchard were then removed and the following year the same Variety 1 trees were examined for affected leaves.

Data:

<u>Year 1</u>		<u>Year 2</u>
Tree	Count	Count
1	38	32
2	10	16
3	84	57
4	36	28
5	50	55
6	35	12
7	73	61
8	48	29

Does the data suggest there is a difference between the mean number of affected leaves in the two periods?

Example: Cedar-apple rust data (cont'd):

Apply 2-sample t-test?

```
> t.test(year1, year2, var.equal=T)
```

Two Sample t-test

data: variety1 and year2

t = 0.9899, df = 14, p-value = 0.3390

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-12.25070 33.25070

sample estimates:

mean of x mean of y

46.75 36.25

### *Paired Samples*

Consider a study comparing weight gains in rats before and after receiving a new dietary regime. Let  $(X_1, Y_1), \dots, (X_n, Y_n)$  be the respective weights of  $n$  rats. Clearly, there is dependence between observations on the same rat, hence, the two-sample procedures discussed above will not be appropriate.

Let  $D_i = X_i - Y_i$ ,  $i = 1, \dots, n$ . Then the problem may be handled based on the  $D_i$ 's using the one-sample t-test (paired t-test) or the Wilcoxon signed-rank test.

# Example: Cedar-apple rust data (cont'd):

```
> t.test(year1,year2,paired=T)
```

Paired t-test

data: variety1 and year2

$t = 2.4342$ ,  $df = 7$ ,

**p-value = 0.04514**

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

**0.2999574 20.7000426**

sample estimates:

mean of the differences

10.5

```
> wilcox.test(year1,year2,paired=T)
```

Wilcoxon signed rank test with continuity correction

data: year1 and year2

$V = 32.5$ ,  $p\text{-value} = 0.04967$

alternative hypothesis: true  $\mu$  is not equal to 0



## Measures of Associations

Given a random sample  $(X_1, Y_1), \dots, (X_n, Y_n)$  on  $n$  individuals, we are often interested in the degree of association or relationship between  $X$  and  $Y$ .

1. *Pearson's product-moment correlation coefficient*

The population correlation coefficient is defined as

$$\rho = \frac{E(X - \mu_x)(Y - \mu_y)}{\sigma_x \sigma_y}$$

and has the property −  $1 \leq \rho \leq 1$ .

The corresponding sample quantity is given by

$$r = \frac{\sum_i (X_i - \bar{X})(Y_i - \bar{Y})}{(n - 1)S_x S_y}$$

with similar properties as  $\rho$ .

For the problem of testing  $H_0 : \rho = 0$  vs.  $H_1 : \rho > 0$ , a suitable test statistic is given by

$$T = \sqrt{\frac{n-2}{1-r^2}}$$

which has a  $t_{n-2}$  null distribution.

The test of the more general hypothesis  $H_0 : \rho = \rho_0$  vs.  $H_1 : \rho > \rho_0$  involves Fisher's  $Z$ -transformation. Let

$$Q(r) = \frac{1}{2} \ln \left( \frac{1+r}{1-r} \right)$$

Then a test statistic is given by

$$Z = \frac{Q(r) - Q(\rho)}{\sqrt{\frac{1}{n-3}}}$$

which has an approximate standard normal null distribution.

For purposes of confidence interval construction, one often uses the following inverse hyperbolic tangent transformation. Let  $H(r) = \operatorname{atanh}(r)$ . Then  $H(r) - H(\rho)$  has an approximate  $N(0, 1/(n - 3))$  distribution. A

100(1 -  $\alpha$ )% confidence interval for  $\rho$  is

$$\operatorname{tanh} \left[ (H(r) \pm Z_{\alpha/2} \frac{1}{\sqrt{n - 3}}) \right].$$

## Remarks

- The product moment correlation coefficient may not be appropriate for ordinal data
- The sample correlation coefficient  $r$  is extremely sensitive to outliers

## 2. *Spearman's rank correlation coefficient*

Spearman's rho is a measure of association based on ranks.

Given  $(X_1, Y_1), \dots, (X_n, Y_n)$ , let  $d_i$  denote the difference between the ranks of  $X_i$  and  $Y_i$ , when the  $X$ 's and  $Y$ 's are ranked separately among themselves.

For tied observation, the mean of the ranks are taken.

When there are no ties, Spearman's rho  $r_s$  coincides with the Pearson product-moment correlation coefficient applied to the ranks, and is given by

$$r_s = 1 - \frac{6 \sum_i d_i^2}{n(n^2 - 1)}$$

### 3. *Kendall's tau*

Kendall's tau is another rank-based measure of association. Let  $R_i$  and  $S_i$  denote the ranks of  $X_i$  and  $Y_i$ , respectively. Then

$$\tau = \frac{1}{\binom{n}{2}} \sum_{i < j} \text{sgn}(R_i - R_j)(S_i - S_j)$$

where

$$\text{sgn}(x) = \begin{cases} 1 & x > 0 \\ 0 & x = 0 \\ -1 & x < 0 \end{cases}$$

```
> cor.test(variety1,year2)
```

Pearson's product-moment  
correlation

data: variety1 and year2

t = 3.973, df = 6,

p-value = 0.007342

alternative hypothesis:

true correlation is not equal to 0

95 percent confidence interval:

0.3662168 0.9725356

sample estimates:

cor

0.851221

```
> cor.test(variety1,year2,method="sp")
```

Spearman's rank correlation rho

data: variety1 and year2

p-value = 0.002232

alternative hypothesis:

true rho is not equal to 0

sample estimates:

rho

0.9285714



## Goodness-of-Fit Tests

### *CDF Test*

Let  $X_1, \dots, X_n$  be a random sample from  $F$ . The empirical distribution function,  $F_n$ , is defined as

$$F_n(x) = \frac{\sum_i I[X_i \leq x]}{n}$$

for  $x \in R$ . As an estimator of  $F$ ,  $F_n$  has several desirable properties.

One way of assessing goodness-of-fit is by looking at a plot of  $F_n$  and  $F$ , superimposed. The `R` function `cdf.compare()` provides such plots.

## *Chi-Square Test of Goodness of Fit*

Let  $X_1, \dots, X_n$  be a random sample from  $F$ . Group the observations  $x_1, \dots, x_n$  into  $k$  classes. Let  $m_i$  and  $M_i$  be the observed and expected, under  $F$ , counts in the  $i$ 'th class, respectively. The chi-square test of goodness-of-fit is given by:

$$\chi_G^2 = \sum_i \frac{(m_i - M_i)^2}{M_i}$$

and has an approximate  $\chi_{k-1}^2$  distribution under  $F$ . In `R`, the function `chisq.gof` may be used.

## *The Kolmogorov-Smirnov Test*

Let  $X_1, \dots, X_n$  be iid F. Then the one-sample K-S test is given by

$$T_{KS} = \sup_x | F_n(x) - F(x) |$$

In the case of two samples, let  $X_1, \dots, X_n$  and  $Y_1, \dots, Y_m$  be independent random samples from F and G, respectively. Then the K-S test for comparing the two distributions is given by:

$$T_{KS} = \sup_x | F_n(x) - G_m(x) |$$

In `R`, the procedures are implemented using `ks.gof()`.

## Bootstrap Tests

The bootstrap may be used to construct approximate confidence intervals and test critical points, when the associated distributions are not tractable.

Consider the problem of comparing two means  $\mu_1$  and  $\mu_2$ , when sampling  $X_1, \dots, X_n$  and  $Y_1, \dots, Y_m$  from  $F$  and  $G$ , respectively.

To implement the bootstrap, one draws pseudo-random samples from  $F_n$  and  $G_m$  under  $H_0 : \mu_1 = \mu_2$ .

- Draw with replacement,  $B$  pseudo-samples:  $X_1^*, \dots, X_n^*$  from  $F_n$  and  $Y_1^*, \dots, Y_m^*$  from  $\tilde{G}_m$ , where  $\tilde{G}_m$  is the edf of  $\tilde{Y}_1, \dots, \tilde{Y}_m$ , and

$$\tilde{Y}_i = Y_i + \bar{X} - \bar{Y}$$

For the  $b$ 'th pseudo-sample, compute

$$Z_b^* = \frac{(\bar{X}^* - \bar{Y}^*)}{\sqrt{\frac{s_1^*}{n} + \frac{s_2^*}{m}}}$$

- Calculate the p-value as

$$p_B = \frac{\sum_{b=1}^B I(|Z_b^*| \geq z_{obs})}{B}$$

where  $z_{obs}$  is the observed value of the test statistic based on the original sample.

# ***Bayesian Statistics***

Bayes theorem:

The diagram illustrates Bayes' theorem with the following components and labels:

- Likelihood:** An arrow points to the term  $P(y|\theta)$  in the numerator.
- Prior:** An arrow points to the term  $P(\theta)$  in the numerator.
- Posterior:** An arrow points to the term  $P(\theta|y)$  on the left side of the equation.
- Marginal Likelihood:** The term  $p(y)$  in the denominator is circled in yellow. An arrow points to it from the text "Ensures integration to 1 i.e. Proper distribution properties".

$$P(\theta|y) = \frac{P(y|\theta) \times P(\theta)}{p(y)}$$

Posterior

Ensures integration to 1  
i.e. Proper distribution properties

Posterior Inference  $\propto$  Likelihood  $\times$  Prior

# Prior Distributions

- Non-informative priors
  - Vague priors, reference priors, flat priors
  - Have minimal impact on the posterior distribution, e.g.,
    - if  $\theta$  is normal( $\theta_o, \sigma_o^2$ ), pick large  $\sigma_o^2$
    - Or, if  $0 \leq \theta \leq 1$ , use  $U(0,1)$  as a prior for  $\theta$
- Informative priors
  - Based on expert opinion
  - From previous studies

# Conjugate Prior Distributions

Family	Conjugate Prior for $\theta$
Bin $(n, \theta)$	$\theta \sim \text{Beta}(\alpha, \beta)$
$\mathcal{P}(\theta)$	$\theta \sim \mathcal{G}(\alpha, \beta)$
$\mathcal{N}(\theta, \sigma^2); \sigma^2$ known	$\theta \sim \mathcal{N}(\theta_0, \sigma_0^2)$
$\mathcal{G}(\alpha, \theta); \alpha$ known	$\theta \sim \mathcal{G}(\delta_0, \gamma_0)$
$\mathcal{N}(\alpha, \theta); \alpha$ known	$\theta^{-1} \sim \mathcal{G}(\delta_0, \gamma_0)$

The likelihood, considered in terms of the parameter(s), has a kernel in the same form as the prior distribution.

Exponential families have conjugate priors in general.

Conjugate priors can be either informative or non-informative.



# The Jeffreys' Prior

A general approach: Define the prior based on the Fisher information *I of the* likelihood

$$I(\theta | \mathbf{x}) = -E_x \left( \frac{\partial^2 \ln \ell(\theta | \mathbf{x})}{\partial \theta^2} \right)$$

So that

$$p(\theta) \propto \sqrt{I(\theta | \mathbf{x})}$$

	<b>Frequentist</b>	<b>Bayesian</b>
Prior information	Informally used in design	Used formally to priors
Interpretation of the parameter of interest	A fixed state of nature	An unknown quantity with probability distribution
Approach	“How likely is the data, given a particular value of the parameter?”	Use posterior distribution of the parameter to calculate and use of the posterior distribution in formal decision analysis
Presentation of results	P-values, confidence intervals	Predictive probabilities

# Credible Intervals

*a  $100(1-\alpha)\%$  two-sided Bayesian credible interval for a single parameter  $\theta$  with equal posterior tail probabilities*

$$[\theta_{\alpha/2}, \theta_{1-\alpha/2}]$$

*where the two end points are the  $\alpha/2$  and  $1 - \alpha/2$  quantiles of the posterior distribution of  $\theta$*

# Bayesian vs Frequentist Inference

- Bayesian estimates with flat priors equivalent to ML estimates
- Asymptotically, priors would be irrelevant and Bayesian and frequentist results would converge.
- Frequentist one-tailed tests equivalent to credible intervals. But, Bayesian methods poor for two-tailed tests.
  - E.g., probability of zero in a continuous distribution is zero.
  - In Bayesian paradigm, frequentist p-values violate the likelihood principle.

# Bayes Factors

Bayes Factors: Analogues of likelihood ratio tests.

- Ratio of prior and posterior information used to provide evidence in favor of one model specification vs another.

Wish to compare two competing models,  $M_1$  and  $M_2$ , with respective parameters  $\theta_1$  and  $\theta_2$ . and likelihood specifications :

$$f_1(x | \theta_1) \text{ and } f_2(x | \theta_2)$$

PosteriorOdds = Prior Odds/Data \* Bayes Factor

$$\frac{\pi(M_1 | x)}{\pi(M_2 | x)} = \frac{p(M_1) / p(x)}{p(M_2) / p(x)} \times \frac{\int_{\theta_1} f_1(x | \theta_1) p_1(\theta_1) d\theta_1}{\int_{\theta_2} f_2(x | \theta_2) p_2(\theta_2) d\theta_2}$$

$$\text{Bayes Factor} = B(x) = \frac{\pi(M_1 | x) / p(M_1)}{\pi(M_2 | x) / p(M_2)}$$

NB: Bayes Factors intuitive, but often quite difficult to calculate.

# Rule of Thumb

$$\text{Bayes Factor} = B(x) = \frac{\pi(M_1 | x) / p(M_1)}{\pi(M_2 | x) / p(M_2)}$$

- $B(x) \geq 1$ , evidence in favor of model 1
- $1 > B(x) \geq 10^{-1/2}$ , weak evidence against model 1.
- $10^{-1/2} > B(x) \geq 10^{-1}$ , substantial evidence against model 1.
- $10^{-1} > B(x) \geq 10^{-2}$ , strong evidence against model 1.
- $10^{-2} > B(x)$ , decisive evidence against model 1.

## Problem Set 2

I) Reading Assignment: Chapter 2,3,4. The Statistical Sleuth: A Course in Methods of Data Analysis. Ramsey & Schafer

II) Consider the RatPupWeight (the weight of rat pups) in R library *nlme*:

`data(RatPupWeight, package="nlme")`

1) *Disregarding the effects of all the other variables, determine whether there is a significant difference between mean weight of rat pups in the control and active (low/high combined) treatment groups using each of the following procedures:*

- a) *A parametric procedure*
- b) *A non-parametric procedure*
- c) *A re-sampling procedure*

2) *Discuss the assumption underlying the analyses in (1) above, their validity, and any remedial measures to be taken..*

3) *Determine whether there is a significant association between weight and Litter Size for each of the following:*

- a) *A parametric procedure*
- b) *A non-parametric procedure*

III) *Suppose data are drawn from a normal distribution, with mean  $\theta$  and variance 1. Assume the prior distribution of  $\theta$  is  $\text{normal}(\theta_0, \sigma^2_0)$ .*

1) *Determine the posterior distribution of  $\theta$ , and*

2) *Indicate the form of a 95% credible interval for the posterior mean of*

A useful reference on R may be found at:

<http://heather.cs.ucdavis.edu/~matloff/132/NSPpart.pdf> (courtesy Jon Fintzi)