

Stat 4201 Homework 4

Mengqi Zong < *mz2326@columbia.edu* >

February 19, 2012

Problem 1

- a. The scatter plot is shown in Fig 1.
- b. The output of the simple linear regression model from R is shown below:

Call:

```
lm(formula = cesium$Mushrooms ~ cesium$Soils)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.279	-30.319	-9.483	31.000	54.271

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	16.72569	12.41954	1.347	0.19807
cesium\$Soils	0.09590	0.02993	3.205	0.00591 **

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 36.56 on 15 degrees of freedom

Multiple R-squared: 0.4064, Adjusted R-squared: 0.3668

F-statistic: 10.27 on 1 and 15 DF, p-value: 0.005909

- c. The output of the simple linear regression model excluding sample number 17 from R is shown below:




Figure 1: Problem 1-a: scatterplot

Call:

```
lm(formula = cesium.trim$Mushrooms ~ cesium.trim$Soils)
```

Residuals:

Min	1Q	Median	3Q	Max
-41.658	-13.938	-4.061	9.744	65.908

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	43.87726	12.25571	3.580	0.00301 **
cesium.trim\$Soils	-0.03693	0.04454	-0.829	0.42085

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 27.76 on 14 degrees of freedom

Multiple R-squared: 0.04682, Adjusted R-squared: -0.02126

F-statistic: 0.6877 on 1 and 14 DF, p-value: 0.4208

d. Based on the scatterplot from part a, we can see that example number 17 is very likely to be an outlier since it is far from every other samples. After excluding example number 17, we get the linear model described in part c. Note that this model's R-squared is 0.0468, which is very close to 0. This indicates that the simple linear regression model is poorly fitted. From the coefficients' p-value, we can see that it is very likely that the model is a constant model.

In conclusion, there is no linear relationship between Cesium concentration in soil and Cesium concentration in mushrooms after Chernobyl accident.

e.

About linearity, based on the R-squared from part b, we can see that the simple linear regression model is poorly fitted.

About normality, I use Shapiro-Wilk normality test on the residuals, here is the result from R:

Shapiro-Wilk normality test

data: res.p1

W = 0.9113, p-value = 0.1053

The p-value is 0.1053. This indicates the data is normal.

About homoscedasticity, I use the Residual vs Fitted plot to do the analysis. The plot is shown in Fig 2. Based on the plot, we can see that the data is not homoscedastic.

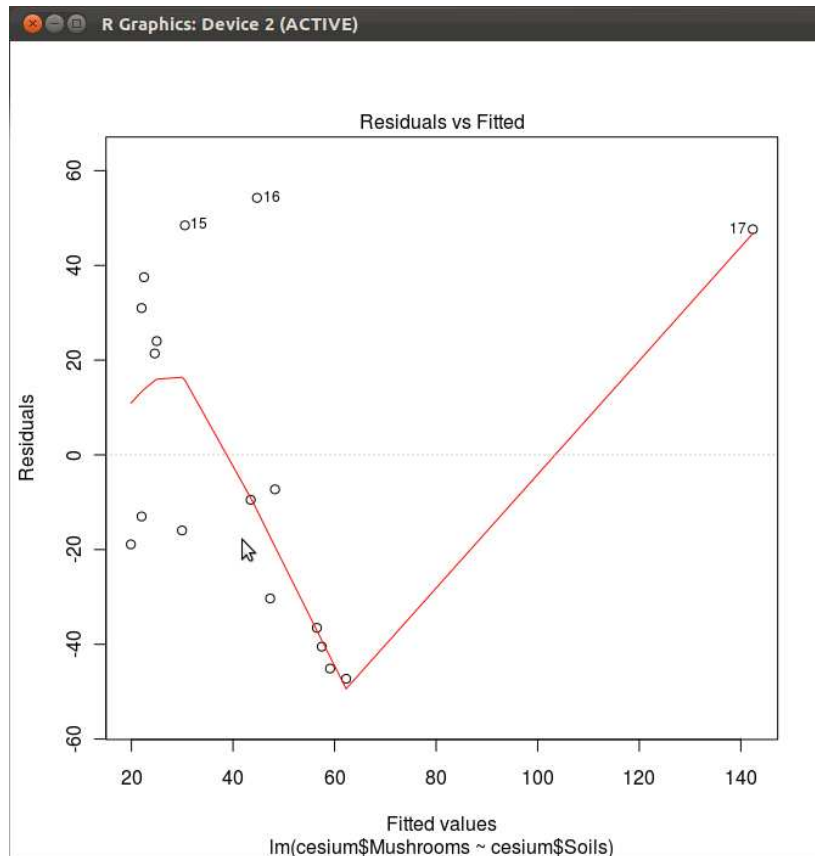


Figure 2: Problem 1-e: Residual vs Fitted

About uncorrelated error, I use Pearson's product-moment test to do the analysis. And here is the result from R:

Pearson's product-moment correlation

```
data: cesium$Mushrooms and cesium$Soils
t = 3.2045, df = 15, p-value = 0.005909
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
```

```

0.2261250 0.8558834
sample estimates:
      cor
0.6374843

```

As we can see, Cesium concentrations in mushrooms and soils are correlated.

Problem 2

a. The output of the simple linear regression model from R is shown below:

```

Call:
lm(formula = stack.loss ~ stack.x[, 1] + stack.x[, 2] + stack.x[,
    3])

```

```

Residuals:
    Min       1Q   Median       3Q      Max
-7.2377 -1.7117 -0.4551  2.3614  5.6978

```

```

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  -39.9197    11.8960  -3.356  0.00375 **
stack.x[, 1]   0.7156     0.1349   5.307  5.8e-05 ***
stack.x[, 2]   1.2953     0.3680   3.520  0.00263 **
stack.x[, 3]  -0.1521     0.1563  -0.973  0.34405
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

```

```

Residual standard error: 3.243 on 17 degrees of freedom
Multiple R-squared:  0.9136, Adjusted R-squared:  0.8983
F-statistic:  59.9 on 3 and 17 DF,  p-value: 3.016e-09

```

b.

About linearity, the R-squared for this model is 0,9136. This indicates that the simple linear regression model is fitted well.

About normality, I use Shapiro-Wilk normality test on the residuals, here is the result from R:

Shapiro-Wilk normality test

```
data: e  
W = 0.974, p-value = 0.8186
```

The p-value is 0.8186. This indicates the data is normal.

About homoscedasticity, I use the Residual vs Fitted plot to do the analysis. The plot is shown in Fig 3. Based on the plot, we can see that the data is nearly homoscedastic.

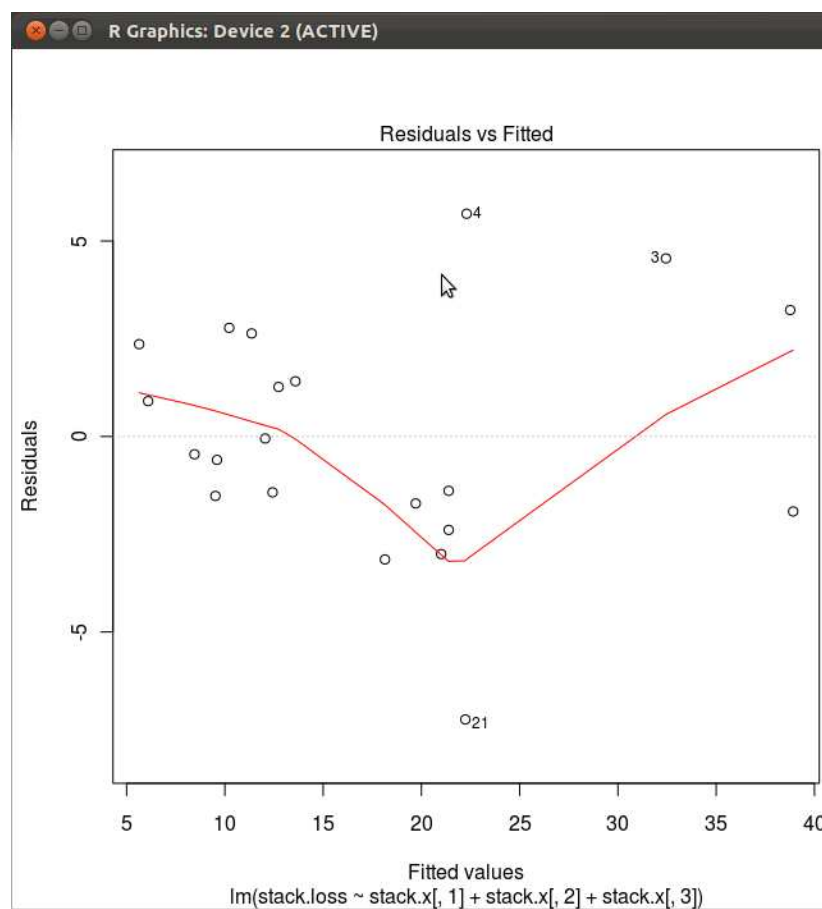


Figure 3: Problem 2-a: Residual vs Fitted

About uncorrelated error, I use Perason's product-moment test to do the analysis. And here is the result from R:

Pearson's product-moment correlation

```
data:  stack.loss and stack.x[, 1]
t = 10.2079, df = 19, p-value = 3.774e-09
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.8092570 0.9673185
sample estimates:
      cor
0.9196635
```

Pearson's product-moment correlation

```
data:  stack.loss and stack.x[, 2]
t = 7.8977, df = 19, p-value = 2.028e-07
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.7134686 0.9486536
sample estimates:
      cor
0.8755044
```

Pearson's product-moment correlation

```
data:  stack.loss and stack.x[, 3]
t = 1.9014, df = 19, p-value = 0.07252
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
-0.03850282 0.70912123
sample estimates:
      cor
0.3998296
```

As we can see, all three explanatory variables – “Air Flow”, “Water Temp” and “Acid Conc” are correlated with “stack.loss”.

About outliers, here is the studentized deleted residuals from R:

1	2	3	4	5	6
---	---	---	---	---	---

1.20947467	-0.70513857	1.61790411	2.05179748	-0.53050364	-0.96320379
7	8	9	10	11	12
-0.82594672	-0.47365206	-1.04858585	0.42618802	0.87829204	0.96670672
13	14	15	16	17	18
-0.46873058	-0.01695002	0.80061639	0.29118502	-0.59958579	-0.14868029
19	20	21			
-0.19719938	0.44311701	-3.33049332			

As we can see, only $|T_{21}| = 3.33 > t_{.975,17} = 2.110$. So only example 21 is an outlier.

About influential points, here is the DFFITS from R:

1	2	3	4	5	6
0.79472051	-0.48132296	0.74415820	0.78788445	-0.12452440	-0.27915632
7	8	9	10	11	12
-0.43767228	-0.25099001	-0.42339897	0.21312348	0.37621097	0.50917672
13	14	15	16	17	18
-0.20268896	-0.00862896	0.38834173	0.11309290	-0.50202110	-0.06503243
19	20	21			
-0.09067758	0.13083298	-2.10029635			

As we can see, only $|DFFITs_{21}| > 1$. So only example 21 is influential.

Appendices

The code is listed below:

```
# problem1
cesium <- read.table("Chernobyl_Fallout", header=TRUE)
#(a)
postscript(file="~/Documents/LaTeX/stat4201-hmwk4/sactterplot.eps",
           onefile=FALSE, horizontal=FALSE)
plot(cesium$Mushrooms~cesium$Soils)
dev.off()

#(b)
cesium.reg <- lm(cesium$Mushrooms~cesium$Soils)
```



```

cesium.reg.summary <- summary(cesium.reg)
print(cesium.reg.summary)
#(c)
cesium.trim <- cesium[c(1:16),]
cesium.trim.reg <- lm(cesium.trim$Mushrooms~cesium.trim$Soils)
cesium.trim.reg.summary <- summary(cesium.trim.reg)
print(cesium.trim.reg.summary)
# additional
lmi.p1 <- lm.influence(cesium.reg)
res.p1 <- resid(cesium.reg)
nt.p1 <- shapiro.test(res.p1)
print(nt.p1)

cor.p1 <- cor.test(cesium$Mushrooms, cesium$Soils)
print(cor.p1)

# problem2

attach(stackloss)
#a)
fit1 <- lm(stack.loss~stack.x[,1] + stack.x[,2] + stack.x[,3])

lmi <- lm.influence(fit1)
lms <- summary(fit1)
print(lms)

#b)
e <- resid(fit1)
nt.p2 <- shapiro.test(e)
print(nt.p2)

cor1.p2 <- cor.test(stack.loss, stack.x[,1])
print(cor1.p2)
cor2.p2 <- cor.test(stack.loss, stack.x[,2])
print(cor2.p2)
cor3.p2 <- cor.test(stack.loss, stack.x[,3])
print(cor3.p2)

```

```
s <- lms$sigma
si <- lmi$sigma
xxi <- diag(lms$cov.unscaled)
h <- lmi$hat

bi <- coef(fit1) - t(coef(lmi))
stand.resid <- e/(s*(1-h)^0.5)
student.resid <- e/(si*(1-h)^0.5)
DFFITS <- h^0.5*e/(si*(1-h))
```